


What's So Funny? - Humor Text Classifier



Erin Cox

<https://github.com/coxem14/Capstone-2>


Background



Chatbots are helping many organizations achieve a higher level of customer service, and now many consumers have virtual assistants like Alexa in their homes, and Google Assistant and Siri on their phones.

Intent detection, like humor detection, is critical to improving and expanding the applications of natural language processing.

Data



The dataset was created by Issa Annamoradnejad and team at Cornell University and was used in the paper *ColBERT: Using BERT Sentence Embedding for Humor Detection*. I downloaded the dataset from Kaggle.

The dataset includes 200k short text samples labeled as either True or False for humor.

EDA - Initial Inspection


	text	humor
0	Joe Biden rules out 2020 bid: 'guys, I'm not running'	False
1	Watch: Darvish gave hitter whiplash with slow pitch	False
2	What do you call a turtle without its shell? dead.	True
3	5 reasons the 2016 election feels so personal	False
4	Pasco police shot Mexican migrant from behind, new autopsy shows	False
5	Martha Stewart tweets hideous food photo, Twitter responds accordingly	False
6	What is a Pokemon master's favorite kind of pasta? Wartortellini!	True

```
Data columns (total 2 columns):  
#      Column  Non-Null Count  Dtype  
---  -  
0     text     200000 non-null  object  
1     humor     200000 non-null  bool
```

```
True      100000  
False     100000  
Name: humor, dtype: int64
```

[illegible]

EDA - Word Clouds



Humorous texts patterns seem to be:

- What did one _ say to the other _?
- What do you call a _ that _?
- What's the difference between _ and _?
- A _ walks into a bar
- 'Knock Knock' jokes

Top 10 words in humorous texts:

- call, say, one, know, go, make, what', joke, peopl, want

EDA - Word Clouds



Train, Test, Split

```
X = texts['text']  
y = texts['humor']  
  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,  
                                                    shuffle=True,  
                                                    stratify=y)
```

The resulting training datasets had 150,000 texts, while the testing datasets had 50,000 texts.

Model Selection

The primary models I built are as follows:

- Multinomial Naive Bayes
- Bernoulli Naive Bayes

I also explored:

- Random Forest Classifier
- Multilayer Perceptron
- KMeans Clustering - Appendix



Pipelines

```
mnb_pipeline = Pipeline([('tfidf', TfidfVectorizer()),  
                           ('model', MultinomialNB())])
```


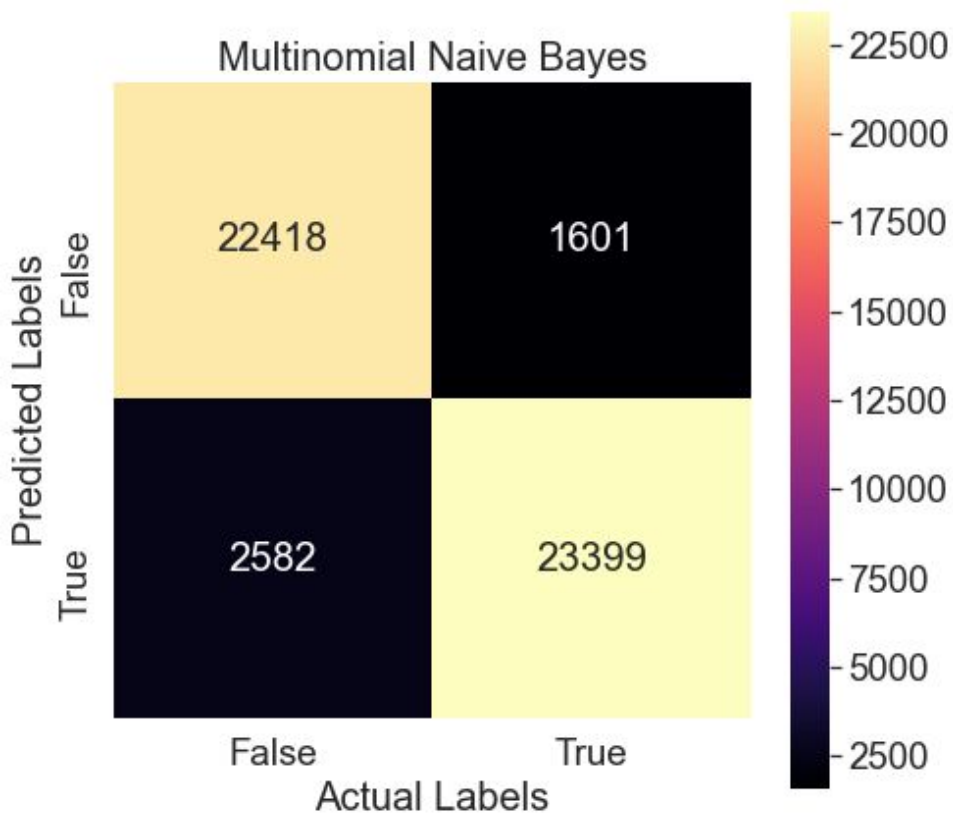
Fitting the Naive Bayes Models

I wanted to see how the models performed with cleaned and uncleaned data inputs, so I ran `X_train` and `X_test` through the corpus cleaner function (same as I used for word clouds) prior to fitting and predicting, respectively.

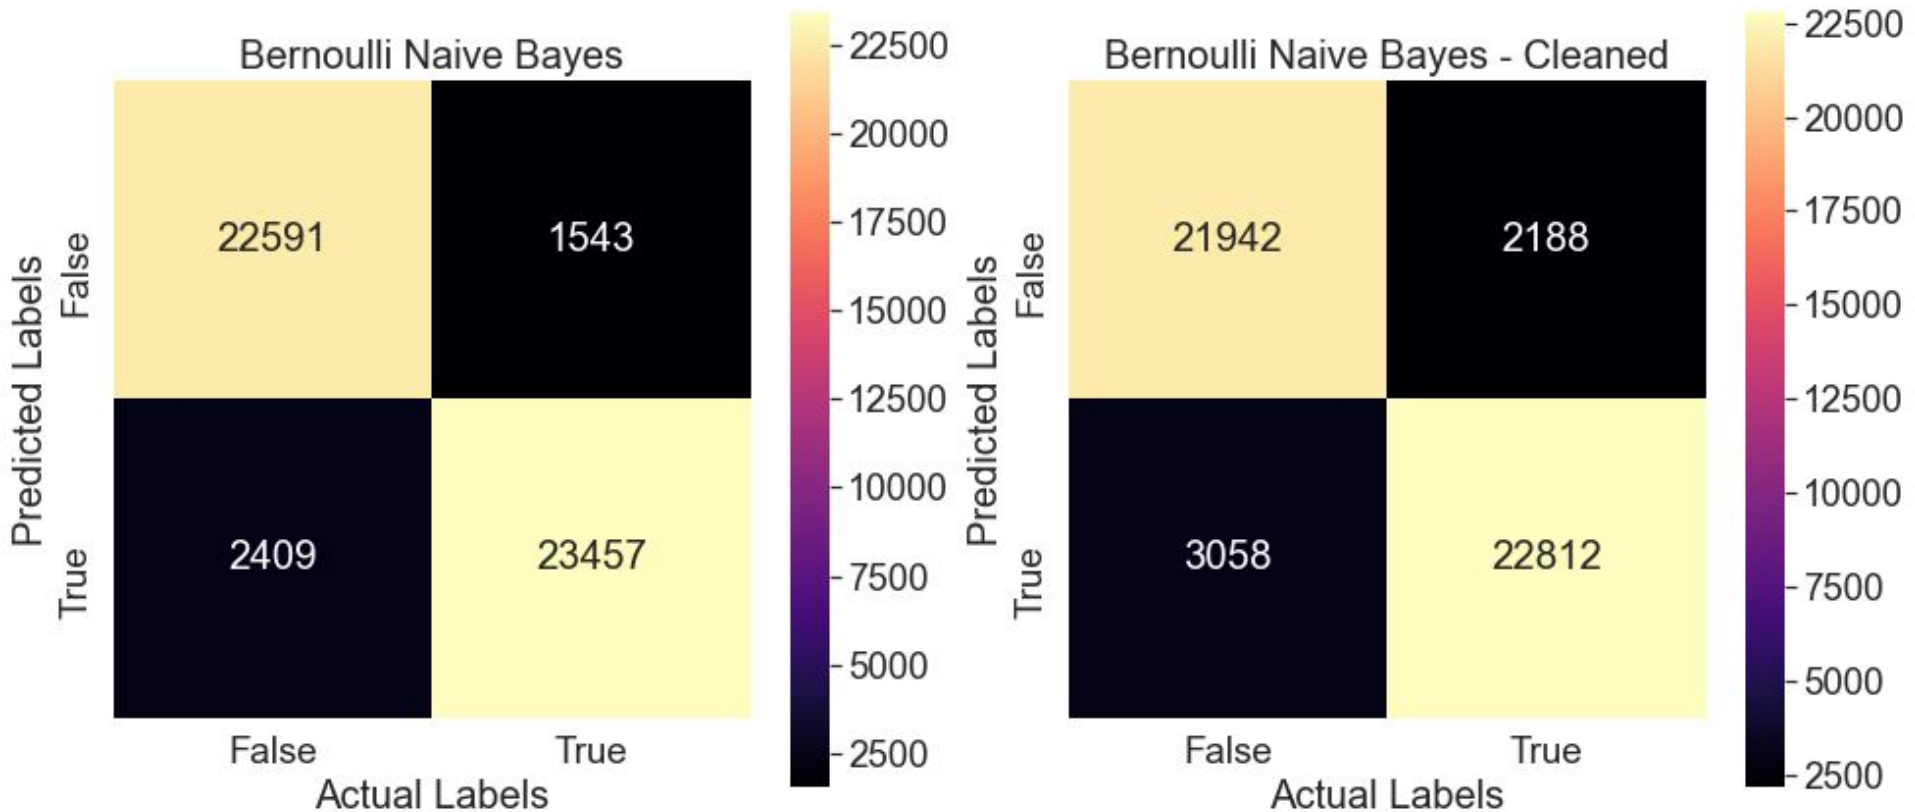
Evaluating the Models

To determine which model performed the best, I fit each model with `X_train` and `y_train`, got the predictions, and compared accuracy, precision, recall scores, confusion matrices, and ROC plots.

Multinomial Naive Bayes



Bernoulli Naive Bayes



Accuracy, Precision, and Recall Scores

Model: Multinomial Naive Bayes

The accuracy on the test set is 0.916.

The precision on the test set is 0.901.

The recall on the test set is 0.936.

Model: Multinomial Naive Bayes – Cleaned

The accuracy on the test set is 0.893.

The precision on the test set is 0.890.

The recall on the test set is 0.896.

Model: Bernoulli Naive Bayes

The accuracy on the test set is 0.921.

The precision on the test set is 0.907.

The recall on the test set is 0.938.

Model: Bernoulli Naive Bayes – Cleaned

The accuracy on the test set is 0.895.

The precision on the test set is 0.882.


The recall on the test set is 0.912.

The model with the highest accuracy: Bernoulli Naive Bayes

The model with the highest precision: Bernoulli Naive Bayes

The model with the highest recall: Bernoulli Naive Bayes

Naive Bayes ROC Plot



Naive Bayes ROC AUC Scores

Model: Multinomial Naive Bayes

The ROC AUC score for the model is 0.975.


Model: Bernoulli Naive Bayes

The ROC AUC score for the model is 0.977.


The model with the largest AUC: Bernoulli Naive Bayes

Random Forest Classifier - Tuning Parameters


RF Accuracy by Max Features



RF Accuracy by Number of Trees



RF Accuracy by Max Depth



```
random_grid = {'n_estimators': n_estimators,
               'max_features': max_features,
               'max_depth': max_depth}

pprint(random_grid)

{'max_depth': [10, 20, 30, 40, 50, 60, 70, 80, 90, 100, 110],
 'max_features': ['auto'],
 'n_estimators': [200, 400, 600, 800, 1000, 1200, 1400, 1600, 1800, 2000]}

rf = RandomForestClassifier()
rf_random = RandomizedSearchCV(estimator = rf,
                              param_distributions = random_grid,
                              n_iter = 3,
                              cv = 3,
                              verbose = 2,
                              n_jobs = -1)

rf_random.fit(X_train_tfidf, y_train)
rf_random.best_params_

{'n_estimators': 1600, 'max_features': 'auto', 'max_depth': 90}
```

The accuracy on the test set is 0.851.
The precision on the test set is 0.888.
The recall on the test set is 0.804.

The accuracy on the test set is 0.907.
The precision on the test set is 0.905.
The recall on the test set is 0.909.

The ROC AUC score for the model is 0.968.

Multilayer Perceptron

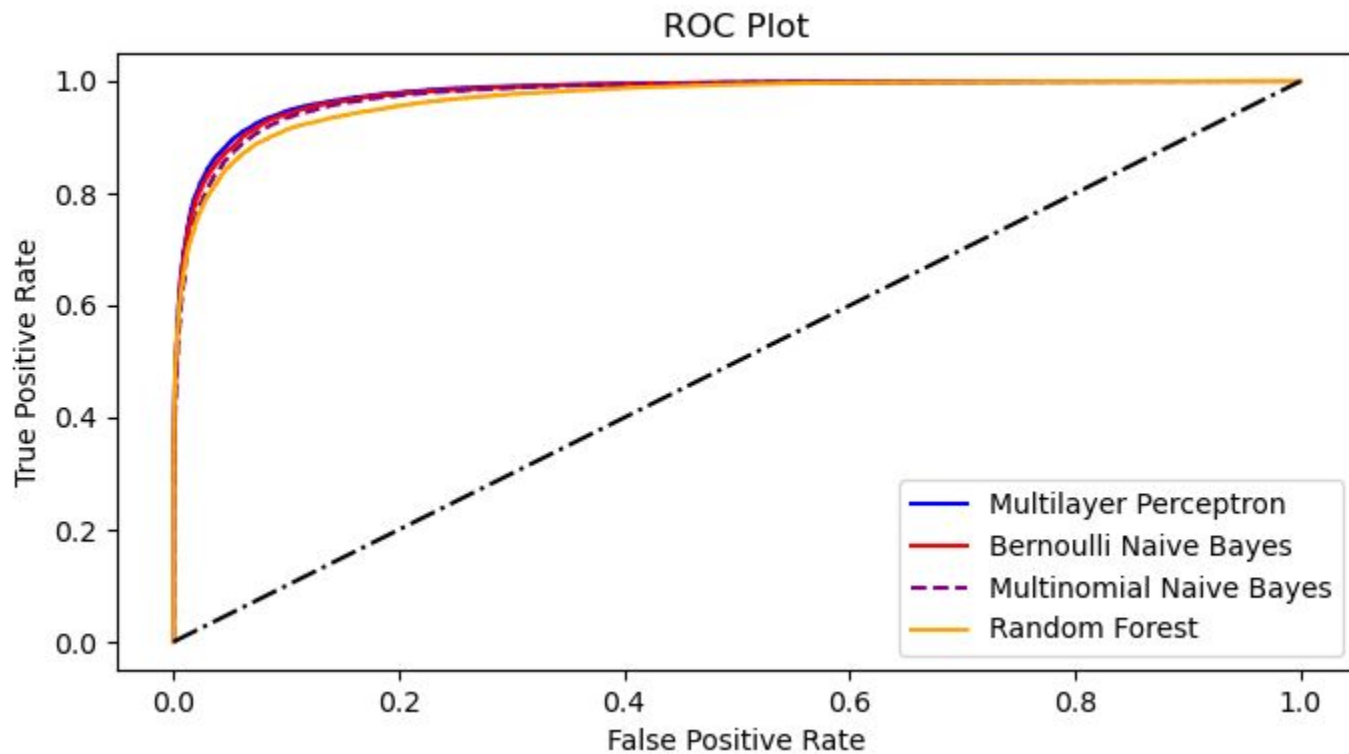
```
mlp_pipeline = Pipeline([('tfidf', TfidfVectorizer(max_features=5000)),  
                          ('model', MLPClassifier(batch_size=32,  
                                                  hidden_layer_sizes=(32,),  
                                                  early_stopping=True))])
```

The accuracy on the test set is 0.925.

The precision on the test set is 0.921.

The recall on the test set is 0.930.

Combined ROC



Problem Words

Top Words in Misclassified Texts



Top 10 words in misclassified texts: say, make, want, people, go, one, kid, day, know, thing

Next Steps

Recap of Findings:

- MLP and Bernoulli Naive Bayes Models performed the best at classifying the texts and predicting on unseen data.
- Common misclassified word included many stopwords, and words which are equally likely to be present in humorous and serious texts.

Future Steps:

- Further cleaning for word clouds
- Further tuning to improve models
- Train, test, split - try cross validation, changing test size
- Featurizing Text - explore n-grams (or sequence of words)
- Random Forest - more in depth hyperparameter tuning
- MLP - more in depth hyperparameter tuning
- Test multiple ks for KMeans
- Explore PCA
- Explore CNN/RNN

Appendix

Results from Cohort Submissions

MLP

Text: When the prosecuting attorney asked Gucci Mane if he was guilty he said, 'bitch I might be.'
Prediction: Humorous
Label: Humorous

Text: Camus says that when you wish yourself into the future you commit suicide by wishing yourself closer to your death.
Prediction: Serious
Label: Serious

Text: How did the random variable get into the club? By showing a fake i.i.d.
Prediction: Humorous
Label: Humorous

Text: Laugh it up! Humor is universal across human cultures – and fuels psychological research on everything from social perception to emotion
Prediction: Humorous
Label: Serious

Text: Saddest 6 word story: "Soup of the Day: No Soup."
Prediction: Serious
Label: Humorous

Text: My mom always told me I wouldn't accomplish anything by lying in bed all day. But look at me now, ma! I'm saving the world!
Prediction: Humorous
Label: Humorous

Text: If I keep stress-eating at this level, the buttons on my shirt will start socially distancing from each other.
Prediction: Humorous
Label: Humorous

Text: To help prevent the spread of COVID-19, everyone should wear a mask in public.
Prediction: Serious
Label: Serious

Text: Avoid close contact with people who are sick.
Prediction: Humorous
Label: Serious

Text: What did one support vector say to another support vector? I feel so marginalized.
Prediction: Humorous
Label: Humorous

BNB

Text: When the prosecuting attorney asked Gucci Mane if he was guilty he said, 'bitch I might be.'
Prediction: Humorous
Label: Humorous

Text: Camus says that when you wish yourself into the future you commit suicide by wishing yourself closer to your death.
Prediction: Humorous
Label: Serious

Text: How did the random variable get into the club? By showing a fake i.i.d.
Prediction: Humorous
Label: Humorous

Text: Laugh it up! Humor is universal across human cultures – and fuels psychological research on everything from social perception to emotion
Prediction: Serious
Label: Serious

Text: Saddest 6 word story: "Soup of the Day: No Soup."
Prediction: Humorous
Label: Humorous

Text: My mom always told me I wouldn't accomplish anything by lying in bed all day. But look at me now, ma! I'm saving the world!
Prediction: Humorous
Label: Humorous

Text: If I keep stress-eating at this level, the buttons on my shirt will start socially distancing from each other.
Prediction: Humorous
Label: Humorous

Text: To help prevent the spread of COVID-19, everyone should wear a mask in public.
Prediction: Serious
Label: Serious

Text: Avoid close contact with people who are sick.
Prediction: Humorous
Label: Serious

Text: What did one support vector say to another support vector? I feel so marginalized.
Prediction: Humorous
Label: Humorous

Combined ROC AUC Scores

Model: Multinomial Naive Bayes

The ROC AUC score for the model is 0.975.

Model: Bernoulli Naive Bayes

The ROC AUC score for the model is 0.977.

Model: Random Forest

The ROC AUC score for the model is 0.968.

Model: Multilayer Perceptron

The ROC AUC score for the model is 0.979.

The model with the largest AUC: Multilayer Perceptron

KMeans Clustering

```
k = 2
vectorizer = TfidfVectorizer(stop_words='english', max_features=1000)
X_vec = vectorizer.fit_transform(X)
features = vectorizer.get_feature_names()
kmeans = KMeans(n_clusters=k, verbose=2)
kmeans.fit(X_vec)

# Find the top 10 features for each cluster.
n_features = 10
top_centroids = kmeans.cluster_centers_.argsort()[:, -1:-(n_features+1):-1]
print("top features (words) for each cluster:")
for num, centroid in enumerate(top_centroids):
    print(f"{num}, {' '.join(features[i] for i in centroid)}")
```

```
top features (words) for each cluster:
0, like, trump, new, just, people, does, don, photos, make, know
1, did, say, hear, cross, road, got, man, chicken, know, guy
```

KMeans Clustering

cluster 0:

Quote: What's a monster's favorite bean? a human bean.

Label: True

Quote: Cheese shop exploded thankfully i was only hit by da brie

Label: True

Quote: Stephanie gilmore's espy awards fashion is spot on (photos)

Label: False

Quote: Why the 2-million pound ready-to-eat chicken recall is extra risky

Label: False

Quote: Worrying is so stupid. it's like carrying an umbrella waiting for it to

Label: True

Quote: Watertown perspective: the boston marathon suspect manhunt on friday

Label: False

cluster 1:

Quote: Did you hear about the guy who stole a dictionary from the library? he

Label: True

Quote: What did the triceratops sit on? its tricerabottom.

Label: True

Quote: What did a bad teacher tell their wisecracking student? don't get smart

Label: True

Quote: What did mozart tell the terminator i'll be bach

Label: True

Quote: Why did the melon plan a big elaborate wedding? because he cantaloupe.

Label: True

Quote: What did iron say to silver after 30 years? you haven't ag-ed a bit.

Label: True

Cluster 0:

False (99587 texts)

True (90012 texts)

Cluster 1:

True (9988 texts)

False (413 texts)

KMeans Silhoutte Plot

For `n_clusters = 2` The average silhouette_score is : 0.01396

